

Indian Institute of Technology Jodhpur
MLOps

Assignment 1
ML Ops

Submitted by: **Jojo Joseph (G23AI2100)**

Google Colab :

https://colab.research.google.com/drive/16uM0n5-4MNq7Ct540Lkxc2pq_mExswpn?usp=sharing

gitHub : https://github.com/TheJojoJoseph/IITJ_MLOpsAssignment1.git

MLOps Assignment

- **Q1. For each of the following scenarios explain the best matching algorithm among :**
 - a. Regression
 - b. Classification
 - c. Clustering
 - d. Association Rule Mining
 - e. Unsupervised Feature Extraction
- **If we form segments among customers with different financial backgrounds, poverty levels, spending habits and income levels**

Ans

CLUSTERING - The goal is to group customers into segments based on their similarities across multiple features (financial background, poverty level, spending habits, income levels). This is an unsupervised task where you don't have labels and want to discover natural groupings.

- **Suppose the camera has to automatically recognize bus from car**

Ans

Classification.Algorithm - This is a supervised learning issue that requires classifying items (bus vs automobile) based on shape, colour, and other attributes.

The system assigns a label (bus or automobile) based on input data.

- Suppose you create a feature in facebook to automatically recognize your friends without labeling

Ans

UNSUPERVISED FEATURE EXTRACTION

This is an example of unsupervised learning, where the system must learn to recognise patterns in data (friend recognition) without pre-labeled data.

- Extract relevant features from user profiles and behaviours without knowing the labels beforehand.

- If we are to identify the accomplices of a thief in a stealing incident based on the call records

Ans

ASSOCIATION RULE MINING

Discovering hidden linkages in a dataset involves identifying patterns or connections between a thief and probable accomplices using call records.

- Association rule mining can find relationships between individuals who communicate often.

- If we are trying to predict future prices of stock based on volume of sales and difference between high and low

Ans

REGRESSION

This is an example of discovering hidden relationships in a dataset, where you want to identify frequent patterns or connections between the thief and potential accomplices based on call records.

- Association rule mining can help identify links between individuals who frequently contact each other.

- Q2. For the given data set boston.csv (for each measure MSE and R Sq) (Submit .ipynb files as attachment and screenshots)

Ans

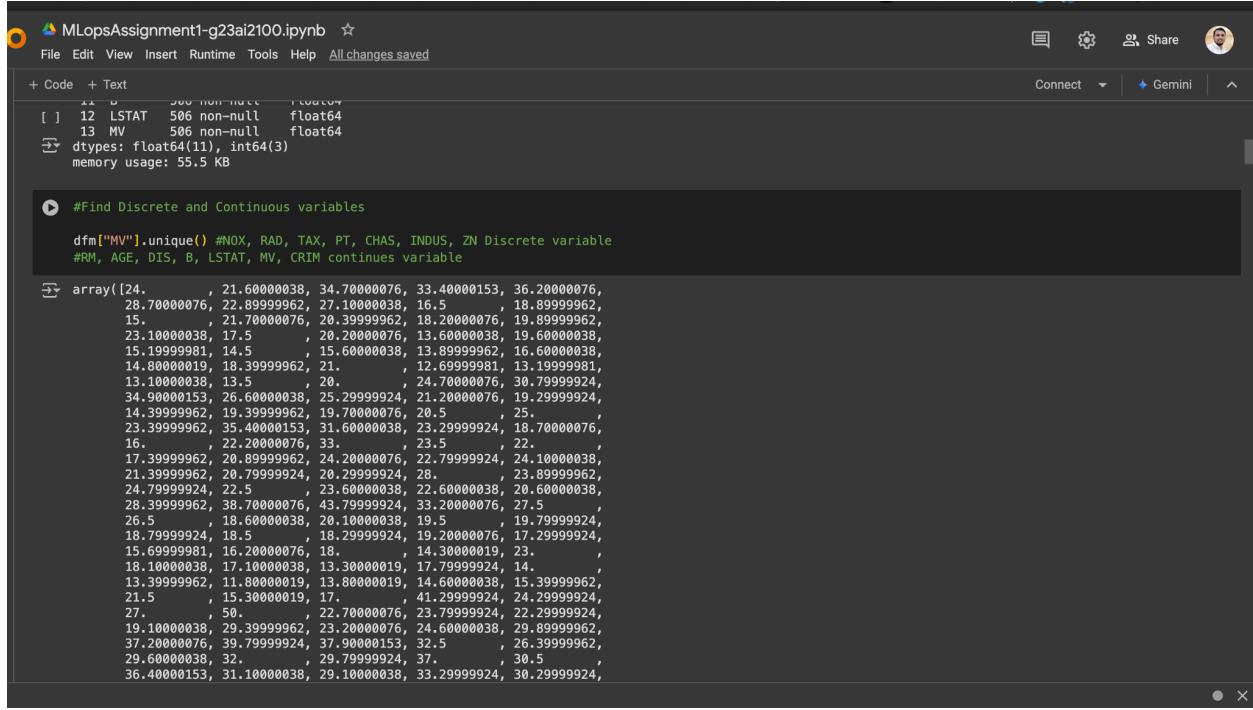
Google Colab :

https://colab.research.google.com/drive/16uM0n5-4MNq7Ct540Lkxc2pq_mExswpn?usp=sharing

GitHub : https://github.com/TheJojoJoseph/IITJ_MLOpsAssignment1.git

- Find Discrete and Continuous variables

Ans



```

MLOpsAssignment1-g23ai2100.ipynb ☆
File Edit View Insert Runtime Tools Help All changes saved
+ Code + Text
[ ] 12 LSTAT 506 non-null float64
13 MV 506 non-null float64
→ dtypes: float64(11), int64(3)
memory usage: 55.5 KB

#Find Discrete and Continuous variables

dfm["MV"].unique() #NOX, RAD, TAX, PT, CHAS, INDUS, ZN Discrete variable
#RM, AGE, DIS, B, LSTAT, MV, CRIM continues variable

→ array([24.          , 21.6000038, 34.7000076, 33.40000153, 36.20000076,
       28.7000076, 22.8999962, 27.1000038, 16.5         ,
       21.7000076, 20.3999962, 18.2000076, 19.8999962,
       23.1000038, 17.5         , 28.2000076, 13.6000038, 19.6000038,
       15.1999981, 14.5         , 15.6000038, 13.8999962, 16.6000038,
       14.8000019, 18.3999962, 21.          , 12.6999981, 13.1999981,
       13.1000038, 13.5         , 20.          , 24.7000076, 30.7999924,
       34.90000153, 26.6000038, 25.2999924, 21.2000076, 19.2999924,
       14.3999962, 19.3999962, 19.7000076, 20.5         ,
       23.3999962, 35.40000153, 31.6000038, 23.2999924, 18.7000076,
       16.          , 22.2000076, 33.          , 23.5         ,
       22.          , 23.5         , 22.          ,
       17.3999962, 20.8999962, 24.2000076, 22.7999924, 24.1000038,
       21.3999962, 20.7999924, 20.2999924, 28.          ,
       23.8999962, 24.7999924, 22.5         , 23.6000038, 22.6000038,
       28.3999962, 38.7000076, 43.7999924, 33.2000076, 27.5         ,
       26.5         , 18.6000038, 20.1000038, 19.5         ,
       18.7999924, 18.5         , 18.2999924, 19.2000076, 17.2999924,
       15.6999981, 16.2000076, 18.          , 14.3000019, 23.          ,
       18.1000038, 17.1000038, 13.3000019, 17.7999924, 14.          ,
       13.3999962, 11.8000019, 13.8000019, 14.6000038, 15.3999962,
       21.5         , 15.3000019, 17.          , 41.2999924, 24.2999924,
       27.          , 50.          , 22.7000076, 23.7999924, 22.2999924,
       19.1000038, 29.3999962, 23.2000076, 24.6000038, 29.8999962,
       37.2000076, 39.7999924, 37.90000153, 32.5         ,
       26.3999962, 29.6000038, 32.          , 29.7999924, 37.          ,
       30.5         ,
       36.40000153, 31.1000038, 29.1000038, 33.2999924, 30.2999924,
       ...
       ]

```

- Find descriptive analytics of all the variables

Ans

```

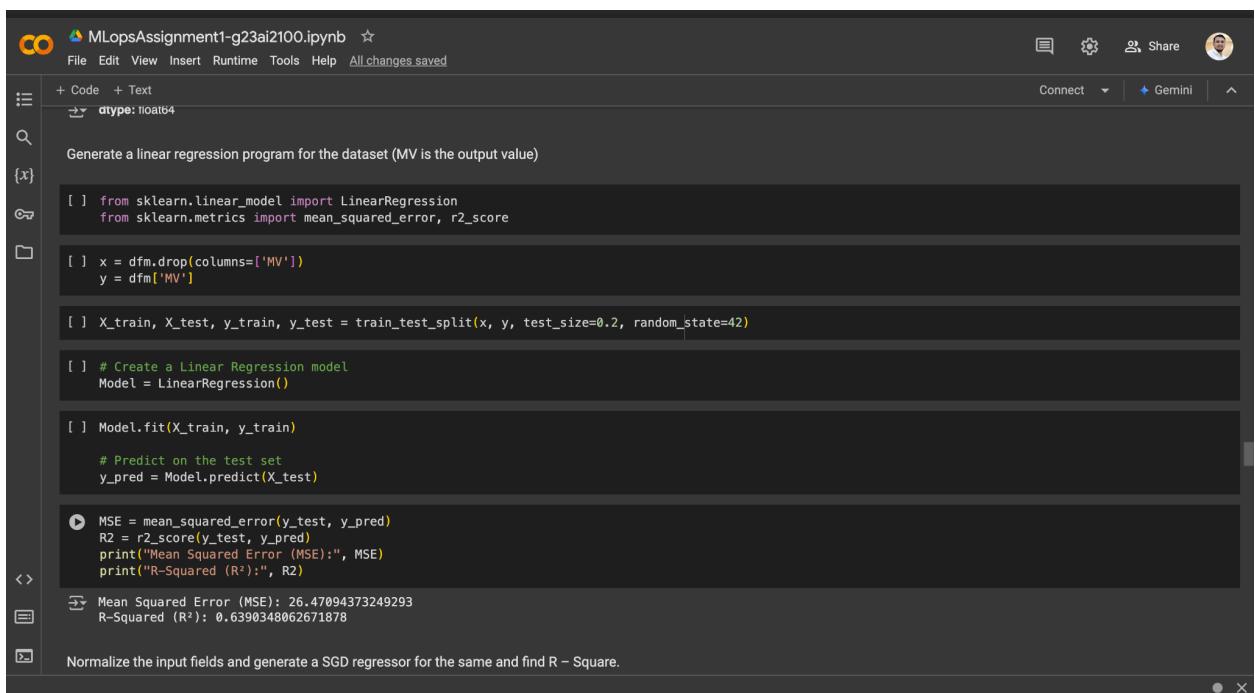
    for clm in dfm.columns:
        print(f"Descriptive Analytics for Column: {clm}")
        print(dfm[clm].describe()) # Descriptive statistics
        print("\n" + "-"*50 + "\n") # Separator for readability

    ⏷ Descriptive Analytics for Column: CRIM
    count    506.000000
    mean     3.613524
    std      8.601545
    min      0.006320
    25%     0.082045
    50%     0.256510
    75%     3.677083
    max     88.976196
    Name: CRIM, dtype: float64
    -----
    Descriptive Analytics for Column: ZN
    count    506.000000
    mean     11.363636
    std      23.322453
    min      0.000000
    25%     0.000000
    50%     0.000000
    75%     12.500000
    max     100.000000
    Name: ZN, dtype: float64
    -----

```

- Generate a linear regression program for the dataset (MV is the output value)

Ans



The screenshot shows a Jupyter Notebook interface with the following details:

- Title:** MLopsAssignment1-g23ai2100.ipynb
- File Menu:** File Edit View Insert Runtime Tools Help All changes saved
- Toolbar:** Code Text Connect Gemini
- Code Cell:**

```

+ Code + Text
↳ ⚡ dtype: float64

Generate a linear regression program for the dataset (MV is the output value)

{x}

[ ] from sklearn.linear_model import LinearRegression
[ ] from sklearn.metrics import mean_squared_error, r2_score

[ ] x = dfm.drop(columns=['MV'])
[ ] y = dfm['MV']

[ ] X_train, X_test, y_train, y_test = train_test_split(x, y, test_size=0.2, random_state=42)

[ ] # Create a Linear Regression model
Model = LinearRegression()

[ ] Model.fit(X_train, y_train)

# Predict on the test set
y_pred = Model.predict(X_test)

[ ] MSE = mean_squared_error(y_test, y_pred)
R2 = r2_score(y_test, y_pred)
print("Mean Squared Error (MSE):", MSE)
print("R-Squared (R²):", R2)

```
- Output Cell:**

```

↳ Mean Squared Error (MSE): 26.47094373249293
R-Squared (R²): 0.6390348062671878

```
- Text Cell:**

```

Normalize the input fields and generate a SGD regressor for the same and find R – Square.

```

- Normalize the input fields and generate a SGD regressor for the same and find R – Square.

```
R-Squared (R2): 0.6361559697887642

Normalize the input fields and generate a SGD regressor for the same and find R – Square.

[ ] from sklearn.preprocessing import StandardScaler
    from sklearn.linear_model import SGDRegressor
    from sklearn.metrics import mean_squared_error, r2_score

[ ] x = dfm.drop(columns=['MV'])
    y = dfm['MV']
    # Normalize the input features
    Scaler = StandardScaler()
    X_scaled = Scaler.fit_transform(x)
    X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.2, random_state=42)

[ ] sgd_regressor = SGDRegressor(max_iter=1000, tol=1e-3, random_state=42)
    sgd_regressor.fit(X_train, y_train)
    y_pred = sgd_regressor.predict(X_test)

[ ] MSE = mean_squared_error(y_test, y_pred)
    R2 = r2_score(y_test, y_pred)

    # Display results
    print("Mean Squared Error (MSE):", MSE)
    print("R-Squared (R2):", R2)

→ Mean Squared Error (MSE): 26.68205970643862
R-Squared (R2): 0.6361559697887642
```

- **Do feature selection by dropping pairwise high correlation fields and run the linear regression and SGD regression after removing irrelevant fields and dropping high correlated fields**

Ans

```

# Linear Regression
linreg = LinearRegression()
linreg.fit(X_train_pca, y_train)
y_pred_linreg = linreg.predict(X_test_pca)
linreg_mse.append(mean_squared_error(y_test, y_pred_linreg))
linreg_r2.append(r2_score(y_test, y_pred_linreg))

# SGD Regressor
sgd = SGDRegressor(max_iter=1000, tol=1e-3, random_state=42)
sgd.fit(X_train_pca, y_train)
y_pred_sgd = sgd.predict(X_test_pca)
sgd_mse.append(mean_squared_error(y_test, y_pred_sgd))
sgd_r2.append(r2_score(y_test, y_pred_sgd))

# Print the results
print("\nResults with Varying Number of PCA Components:\n")
print("Components\tLinear Regression MSE\tLinear Regression R2\tSGD MSE\tSGD R2")
for i, n in enumerate(Components_list):
    print(f"\t{n}\t{linreg_mse[i]:.4f}\t{linreg_r2[i]:.4f}\t{sgd_mse[i]:.4f}\t{sgd_r2[i]:.4f}")

```

Results with Varying Number of PCA Components:

Components	Linear Regression MSE	Linear Regression R2	SGD MSE	SGD R2
1	39.9893	0.4547	39.9882	0.4547
2	34.1245	0.5347	33.9756	0.5367
3	29.5241	0.5974	29.4339	0.5986
4	28.2844	0.6143	28.3277	0.6137
5	30.0843	0.5898	30.1588	0.5887
6	29.8188	0.5934	29.8853	0.5925
7	29.8093	0.5935	29.8798	0.5926
8	29.7089	0.5949	29.7422	0.5944
9	28.5572	0.6106	28.6777	0.6089
10	28.7325	0.6082	28.8298	0.6069
11	27.0729	0.6308	27.2941	0.6278

- Run PCA on the inputs above and perform both SGD and Linear Regression with varying different no of components

Ans

```

[ ] from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score

[ ] x = dfm.drop(columns=['MV'])
y = dfm['MV']

[ ] X_train, X_test, y_train, y_test = train_test_split(x, y, test_size=0.2, random_state=42)

[ ] # Create a Linear Regression model
Model = LinearRegression()

[ ] Model.fit(X_train, y_train)

# Predict on the test set
y_pred = Model.predict(X_test)

[ ] MSE = mean_squared_error(y_test, y_pred)
R2 = r2_score(y_test, y_pred)
print("Mean Squared Error (MSE):", MSE)
print("R-Squared (R²):", R2)

→ Mean Squared Error (MSE): 26.47094373249293
R-Squared (R²): 0.6390348062671878

```

Normalize the input fields and generate a SGD regressor for the same and find R – Square.

```

+ Code + Text
Connect ▾ Gemini ▾
linreg_mse.append(mean_squared_error(y_test, y_pred_linreg))
linreg_r2.append(r2_score(y_test, y_pred_linreg))

# SGD Regressor
sgd = SGDRegressor(max_iter=1000, tol=1e-3, random_state=42)
sgd.fit(X_train_pca, y_train)
y_pred_sgd = sgd.predict(X_test_pca)
sgd_mse.append(mean_squared_error(y_test, y_pred_sgd))
sgd_r2.append(r2_score(y_test, y_pred_sgd))

# Print the results
print("\nResults with Varying Number of PCA Components:")
print("Components\tLinear Regression MSE\tLinear Regression R2\tSGD MSE\tSGD R2")
for i, n in enumerate(Components_List):
    print(f'{n}\t{linreg_mse[i]:.4f}\t{linreg_r2[i]:.4f}\t{sgd_mse[i]:.4f}\t{sgd_r2[i]:.4f}')


→ Results with Varying Number of PCA Components:
Components      Linear Regression MSE      Linear Regression R2      SGD MSE      SGD R2
1              39.9893                  0.4547      39.9882      0.4547
2              34.1245                  0.5347      33.9756      0.5367
3              29.5241                  0.5974      29.4339      0.5986
4              28.2844                  0.6143      28.3277      0.6137
5              30.0843                  0.5808      30.1588      0.5887
6              29.8188                  0.5934      29.8853      0.5925
7              29.8093                  0.5935      29.8798      0.5926
8              29.7089                  0.5949      29.7422      0.5944
9              28.5572                  0.6106      28.6777      0.6089
10             28.7325                  0.6082      28.8298      0.6069
11             27.0729                  0.6308      27.2941      0.6278
12             26.4709                  0.6390      26.7404      0.6354

```

- Perform one hot encoding of discrete variables and perform linear regression

Ans

```
+ Code + text
↳
encoded_feature_names = encoder.get_feature_names_out(discrete_cols)

# Combine the encoded features with the continuous ones
continuous_cols = dfm.drop(columns=discrete_cols + ["MV"], axis=1).columns
data_encoded = pds.DataFrame(encoded_features, columns=encoded_feature_names)
data_continuous = dfm[continuous_cols].reset_index(drop=True)
X = pds.concat([data_continuous, data_encoded], axis=1)
y = dfm["MV"]

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Standardize the continuous features
scaler = StandardScaler()
X_train[continuous_cols] = scaler.fit_transform(X_train[continuous_cols])
X_test[continuous_cols] = scaler.transform(X_test[continuous_cols])

# Train Linear Regression
linreg = LinearRegression()
linreg.fit(X_train, y_train)

# Predict and Evaluate
y_pred = linreg.predict(X_test)
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)

# Print Results
print(f"Mean Squared Error: {mse:.4f}")
print(f"R-Square: {r2:.4f}")

<>
Mean Squared Error: 22.3439
R-Square: 0.6953
```

- Do a optuna based hyperparameter search for KNN regression, Random Forest Regression and Ridge regression for the same data and check the best hyperaparameter values

Ans

```
import optuna

from sklearn.model_selection import train_test_split, cross_val_score

from sklearn.neighbors import KNeighborsRegressor

from sklearn.ensemble import RandomForestRegressor

from sklearn.linear_model import Ridge

from sklearn.preprocessing import StandardScaler

from sklearn.pipeline import Pipeline

from sklearn.metrics import mean_squared_error

import pandas as pds

import numpy as npy
```

```
x = dfm.drop(columns=["MV"], axis=1)

y = dfm["MV"]

# Split the dataset

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
random_state=42)

# Standardization pipeline

Scaler = StandardScaler()

X_train = Scaler.fit_transform(X_train)

X_test = Scaler.transform(X_test)

# Define objective functions for each model

def knn_objective(trial):

    n_neighbors = trial.suggest_int("n_neighbors", 1, 20)

    weights = trial.suggest_categorical("weights", ["uniform", "distance"])

    p = trial.suggest_int("p", 1, 2)

    model = KNeighborsRegressor(n_neighbors=n_neighbors, weights=weights,
p=p)
```

```

    score = cross_val_score(model, X_train, y_train,
scoring="neg_mean_squared_error", cv=5).mean()

    return -score


def rf_objective(trial):

    n_estimators = trial.suggest_int("n_estimators", 10, 200)

    max_depth = trial.suggest_int("max_depth", 3, 20)

    min_samples_split = trial.suggest_int("min_samples_split", 2, 10)

    model = RandomForestRegressor(n_estimators=n_estimators,
max_depth=max_depth, min_samples_split=min_samples_split, random_state=42)

    score = cross_val_score(model, X_train, y_train,
scoring="neg_mean_squared_error", cv=5).mean()

    return -score


def ridge_objective(trial):

    alpha = trial.suggest_loguniform("alpha", 1e-3, 1e2)

    model = Ridge(alpha=alpha)

    score = cross_val_score(model, X_train, y_train,
scoring="neg_mean_squared_error", cv=5).mean()

    return -score


# Run Optuna for each model

print("Optimizing KNN...")

```

```
study_knn = optuna.create_study(direction="minimize")

study_knn.optimize(knn_objective, n_trials=50)

print("Optimizing Random Forest...")

study_rf = optuna.create_study(direction="minimize")

study_rf.optimize(rf_objective, n_trials=50)

print("Optimizing Ridge...")

study_ridge = optuna.create_study(direction="minimize")

study_ridge.optimize(ridge_objective, n_trials=50)

# Print the best parameters and scores

print("Best parameters for KNN:", study_knn.best_params)

print("Best score for KNN:", study_knn.best_value)

print("Best parameters for Random Forest:", study_rf.best_params)

print("Best score for Random Forest:", study_rf.best_value)

print("Best parameters for Ridge:", study_ridge.best_params)

print("Best score for Ridge:", study_ridge.best_value)
```

```
+ Code + Text
    print("Best parameters for KNN:", study_knn.best_params)
    print("Best score for KNN:", study_knn.best_value)

    print("Best parameters for Random Forest:", study_rf.best_params)
    print("Best score for Random Forest:", study_rf.best_value)

    print("Best parameters for Ridge:", study_ridge.best_params)
    print("Best score for Ridge:", study_ridge.best_value)

    e: 14.890279234297765 and parameters: {'n_estimators': 153, 'max_depth': 13, 'min_samples_split': 2}. Best is trial 27 with value: 14.890279234297765.
e: 15.034831763620952 and parameters: {'n_estimators': 129, 'max_depth': 13, 'min_samples_split': 2}. Best is trial 27 with value: 14.890279234297765.
e: 15.120791059579593 and parameters: {'n_estimators': 111, 'max_depth': 13, 'min_samples_split': 5}. Best is trial 27 with value: 14.890279234297765.
e: 14.948749338524511 and parameters: {'n_estimators': 200, 'max_depth': 13, 'min_samples_split': 2}. Best is trial 27 with value: 14.890279234297765.
e: 14.93468054768377 and parameters: {'n_estimators': 199, 'max_depth': 13, 'min_samples_split': 2}. Best is trial 27 with value: 14.890279234297765.
e: 14.948749338524511 and parameters: {'n_estimators': 200, 'max_depth': 13, 'min_samples_split': 2}. Best is trial 27 with value: 14.890279234297765.
e: 15.119225230207885 and parameters: {'n_estimators': 177, 'max_depth': 15, 'min_samples_split': 3}. Best is trial 27 with value: 14.890279234297765.
e: 15.12131980386456 and parameters: {'n_estimators': 198, 'max_depth': 12, 'min_samples_split': 4}. Best is trial 27 with value: 14.890279234297765.
e: 15.102460927871766 and parameters: {'n_estimators': 173, 'max_depth': 15, 'min_samples_split': 2}. Best is trial 27 with value: 14.890279234297765.
e: 15.173784056413613 and parameters: {'n_estimators': 190, 'max_depth': 13, 'min_samples_split': 3}. Best is trial 27 with value: 14.890279234297765.
e: 15.379926814049913 and parameters: {'n_estimators': 183, 'max_depth': 11, 'min_samples_split': 5}. Best is trial 27 with value: 14.890279234297765.
e: 14.929043275526762 and parameters: {'n_estimators': 148, 'max_depth': 16, 'min_samples_split': 2}. Best is trial 27 with value: 14.890279234297765.
e: 15.052433379536364 and parameters: {'n_estimators': 146, 'max_depth': 19, 'min_samples_split': 3}. Best is trial 27 with value: 14.890279234297765.
e: 14.981375537505759 and parameters: {'n_estimators': 116, 'max_depth': 17, 'min_samples_split': 4}. Best is trial 27 with value: 14.890279234297765.
e: 15.100226638189179 and parameters: {'n_estimators': 185, 'max_depth': 16, 'min_samples_split': 2}. Best is trial 27 with value: 14.890279234297765.
e: 15.127046808797544 and parameters: {'n_estimators': 166, 'max_depth': 14, 'min_samples_split': 2}. Best is trial 27 with value: 14.890279234297765.
e: 15.076832181971463 and parameters: {'n_estimators': 200, 'max_depth': 12, 'min_samples_split': 3}. Best is trial 27 with value: 14.890279234297765.
e: 15.11916378253676 and parameters: {'n_estimators': 176, 'max_depth': 17, 'min_samples_split': 2}. Best is trial 27 with value: 14.890279234297765.
e: 15.008700388933104 and parameters: {'n_estimators': 141, 'max_depth': 15, 'min_samples_split': 2}. Best is trial 27 with value: 14.890279234297765.
e: 14.972248305795087 and parameters: {'n_estimators': 101, 'max_depth': 13, 'min_samples_split': 3}. Best is trial 27 with value: 14.890279234297765.
e: 15.008513301054165 and parameters: {'n_estimators': 78, 'max_depth': 16, 'min_samples_split': 2}. Best is trial 27 with value: 14.890279234297765.
e: 15.247876181181431 and parameters: {'n_estimators': 189, 'max_depth': 14, 'min_samples_split': 4}. Best is trial 27 with value: 14.890279234297765.
e: 15.017555034057612 and parameters: {'n_estimators': 152, 'max_depth': 12, 'min_samples_split': 3}. Best is trial 27 with value: 14.890279234297765.
```

- Develop a small web page in flask based on input data and predicting MV output with different input data given

app.py

TechStack : flask

```
from flask import Flask, request, render_template
import joblib
import pandas as pd
app = Flask(__name__)
model = joblib.load('linear_regression_model.pkl')
scaler = joblib.load('scaler.pkl')

feature_names = ['CRIM', 'ZN', 'INDUS', 'CHAS', 'NOX', 'RM',
                 'AGE', 'DIS', 'RAD', 'TAX', 'PTRATIO', 'B', 'LSTAT']

@app.route('/', methods=['GET'])
def home():
    return render_template('index.html', feature_names=feature_names)

# prediction
@app.route('/predict', methods=['POST'])
def predict():
    try:
```

```
# Collect input features from the form
input_features = []
for feature in feature_names:
    value = request.form.get(feature)
    if not value:
        raise ValueError(f"Value for '{feature}' is missing.")
    try:
        input_features.append(float(value))
    except ValueError:
        raise ValueError(f"Invalid value for '{feature}'. Please enter a numeric value.")

input_data = pd.DataFrame(
    [input_features], columns=feature_names) # Convert DataFrame
input_scaled = scaler.transform(input_data) # Scale input features
prediction = model.predict(input_scaled) # Make prediction
prediction = round(prediction[0], 2)
return render_template('result.html', prediction=prediction)

except Exception as e:
    return render_template('index.html', error=str(e), feature_names=feature_names)

if __name__ == '__main__':
    app.run(debug=True)
```



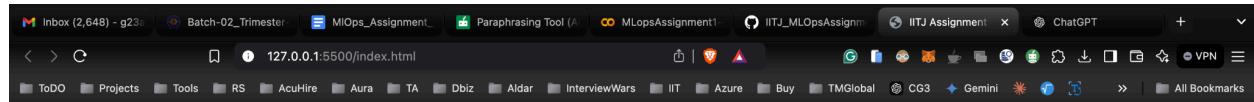
```

# Training
model = LinearRegression()
model.fit(X_train_scaled, y_train)
# Save the scaler and model
joblib.dump(scaler, 'scaler.pkl')
joblib.dump(model, 'linear_regression_model.pkl')

print("Model and scaler complete")

```

UI and index.html



Predict MV Value

CRIM:

ZN:

INDUS:

CHAS:

NOX:

RM:

AGE:

DIS:

RAD:

```

<!DOCTYPE html>
<html lang="en">
  <head>
    <meta charset="UTF-8" />
    <meta name="viewport" content="width=device-width, initial-scale=1.0" />
    <title>IITJ Assignment 1 MLops - g23ai2100</title>
    <style>
      body {
        font-family: Arial, sans-serif;
        line-height: 1.6;
      }
    </style>
  </head>
  <body>
    <h3>Predict MV Value</h3>
    <form>
      <p>CRIM:<br/><input type="text"/></p>
      <p>ZN:<br/><input type="text"/></p>
      <p>INDUS:<br/><input type="text"/></p>
      <p>CHAS:<br/><input type="text"/></p>
      <p>NOX:<br/><input type="text"/></p>
      <p>RM:<br/><input type="text"/></p>
      <p>AGE:<br/><input type="text"/></p>
      <p>DIS:<br/><input type="text"/></p>
      <p>RAD:<br/><input type="text"/></p>
    </form>
  </body>
</html>

```

```
        margin: 20px;
    }
    form {
        max-width: 400px;
        margin: 0 auto;
    }
    label {
        display: block;
        margin: 10px 0 5px;
    }
    input {
        width: 100%;
        padding: 8px;
        margin-bottom: 10px;
        border: 1px solid #ccc;
        border-radius: 4px;
    }
    button {
        background-color: #007bff;
        color: white;
        border: none;
        padding: 10px 15px;
        border-radius: 4px;
        cursor: pointer;
    }
    button:hover {
        background-color: #0056b3;
    }
    h2 {
        text-align: center;
    }
    h3 {
        text-align: center;
        color: green;
    }
</style>
</head>
<body>
    <h2>Predict MV Value</h2>
    <form method="post" action="/">
        <label for="CRIM">CRIM:</label>
        <input type="number" name="CRIM" id="CRIM" required step="any" />
```

```
<label for="ZN">ZN:</label>
<input type="number" name="ZN" id="ZN" required step="any" />

<label for="INDUS">INDUS:</label>
<input type="number" name="INDUS" id="INDUS" required step="any" />

<label for="CHAS">CHAS:</label>
<input type="number" name="CHAS" id="CHAS" required step="any" />

<label for="NOX">NOX:</label>
<input type="number" name="NOX" id="NOX" required step="any" />

<label for="RM">RM:</label>
<input type="number" name="RM" id="RM" required step="any" />

<label for="AGE">AGE:</label>
<input type="number" name="AGE" id="AGE" required step="any" />

<label for="DIS">DIS:</label>
<input type="number" name="DIS" id="DIS" required step="any" />

<label for="RAD">RAD:</label>
<input type="number" name="RAD" id="RAD" required step="any" />

<label for="TAX">TAX:</label>
<input type="number" name="TAX" id="TAX" required step="any" />

<label for="PTRATIO">PTRATIO:</label>
<input type="number" name="PTRATIO" id="PTRATIO" required step="any" />

<label for="B">B:</label>
<input type="number" name="B" id="B" required step="any" />

<label for="LSTAT">LSTAT:</label>
<input type="number" name="LSTAT" id="LSTAT" required step="any" />

<button type="submit">Predict</button>
</form>
<h3>{{ prediction_text }}</h3>
</body>
</html>
```

